



# FUNDAMENTALS OF IMAGE AND VIDEO PROCESSING

Part 8: Image Coding (short version)

# Why coding images


- Visual data are large object to store or transmit
  - Consider an image at HDTV resolution (1920x1080 pixels), 24 bits per pixel (8xRGB): in raw format it will require  $1920 \times 1080 \times 24 = \mathbf{49766400 \text{ bits}}$  (approximately 6Mbytes)
  - If we have a sequence of such images, at 30fps progressive, the bitrate that we get is about **1.5 Gbps** (gigabit per second)
- Big problems both for transmission and storage!
- To allow transmitting an image in realtime or streaming a video we need to reduce the above amounts of information of 1-2 orders magnitude
  - This can be done thanks to **image and video compression**
  - The objective of compression is to obtain media with similar quality using a significantly lower number of bits

# How to compress

- This goal may be achieved by exploiting some intrinsic (and inter-related) characteristics of image and video data:
  - **Low frequency content:** images are a low-pass signals, their frequency content is limited
  - **Slow variation:** image signals are locally stationary, they change slowly along space (and time, for video), except for specific transition points (edges), which are however quite sparse
  - **High-correlation:** if we measure the spatial (temporal) correlation of an image (video) signal, we see that it is a long-memory process
  - **Predictability:** future values of image signals can be often predicted with high accuracy from previous value. This is even more evident for video signals along time
  - **Self-similarity:** parts of the image are sometimes repetitive, even if they show significant variations (think about textures)

# CoDec

- We call CoDec (acronym of Coder+Decoder) a system that transforms an image in a coded (compressed) format and then back to an image
  - The compression module (**encoder**) generates a code that (typically) requires less bits than the original image
  - The decompression module (**decoder**) reconstructs an image (close to, but in general different from the original) from the code



*NB. The code is not an image (pixel), but a set of data that is "meaningful" only for an appropriate decoder*

# CoDec performance

- We typically measure the performance of a CoDec using two parameters:
  - **Compression factor:** it is expressed by **quantitative parameters** that indicate how effective the compression was in terms of bit reduction
  - **Quality:** it is expressed by **quantitative or qualitative parameters** that indicate how accurate is the reconstruction  $I'(x,y)$  with respect to the original image  $I(x,y)$

*NB. While the compression can be measured exactly, the quality is more difficult to evaluate, as it is influenced by subjective factors (human perception)*

# Transform coding


- The term transform coding is associated to a large class of methods that use as a primary tool the transformation of the image in a different domain
  - The transform we have learnt in this course is Fourier, but there are many other transforms that are more suitable for image coding
  - The two most important ones are DCT and DWT (we'll see them)
- Different methods are characterized by:
  - Which image transform they adopt
  - How the transform is applied (whole image, blocks)
  - How the transformed coefficients are coded
- Typically the three things are related each-other



# Transform coding using DCT

- A largely used transform in image coding is DCT
- DCT stands for Discrete Cosine Transform. It is a rather 'young' transform, invented in 1972 by Nasir Ahmed
- DCT is quite similar to DFT but:
  - It is real and not complex (simpler and better for compression)
  - It is more effective in compressing the signal energy in few coefficients
  - It can be calculated with integer computation and it is very fast
- From a mathematical viewpoint, DCT is a transformation using real cosine waves at multiples of a bases frequency along x and y as basis functions
  - In the following slide the basis images of the DCT are shown

## 2D DCT basis functions



8x8 2D-DCT basis functions (64 8x8 images)



## DCT: Example


- ▶ This is how the DCT of Lena appears:
  - ▶ Significant coefficients are very few, concentrated in the top-left corner (the DC coefficient)
  - ▶ The number of coefficients is equal to the image size (as for the DFT)





# Coding of transformed image

- Since the coefficients are sparse and collected into a small portion of the transformed images, there are 2 main ways to encode them efficiently:
  - **Zonal coding:** only the coefficients belonging to the area that most probably will contain significant values are saved, all the others are set to zero
  - **Threshold coding:** only the coefficients that have a value above a given threshold are saved, all the others are set to zero
- Both methods have **pros and cons:**
  - Zonal: does not require storing the position of saved coefficients (more effective in compression) BUT introduces a severe blurring
  - Threshold: more adaptive (better quality) but wastes a lot of bits to store the coefficients' positions


# Zonal vs. threshold coding




coefficients image



a couple of typical zonal masks



coefficients image



selected coefficients and relevant bit-map

```

11001000000000000000
10000000000010000000
00100000000000000000
10000000000000000000
00000010000000000000
10000000000000000000
00000000000000000000
00000000000000000000
00000000000000000000

```

# The JPEG standard

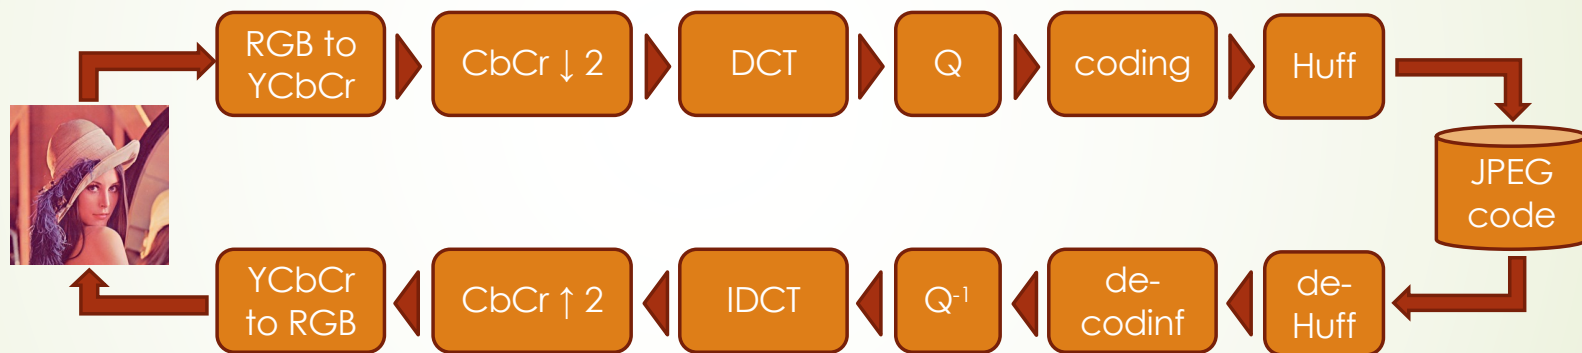
- JPEG is the acronym of **Joint Photographic Experts Group**, a Committee established by ISO (International Standardization Organization) in 1986 and still active
- The first objective of the Committee was to standardize a technique for still picture coding
- The standard known as JPEG (.jpg format) has been released in 1992 and collected the best of the state-of-art worldwide in image compression
- It has been (and still is) extremely successful and widely used, even if several more efficient coders have been made available since then (e.g., JPEG XR, JPEG2000, JPEG XS, etc.)
- It is used on digital cameras, smartphones, social networks, web, professional/amateur repositories, medical archives, etc.

# JPEG standard

- As all standards it is formulated in a very detailed way
  - To ensure interoperability, the format should be very precise
  - The encoder instead, has some degree of freedom
- It includes variants that are rarely implemented (commercial implementations are typically subsets of the standard)
- The target are photographic pictures (photos), but it works also on professional pictures (e.g., biomedical, remote sensing, documents) to some extent
- The target compression is around 1:10-20, with acceptable quality (low artifacts, PSNR > 30 dB)
- It is easy-to-use by design (it has a single parameter)

# JPEG: The algorithm

- We will just focus on the algorithmic part of the standard
  - JPEG is just a sophisticated **block transform coder** using DCT
  - The co-dec scheme is depicted below:




- We will analyze step by step the encoding process
  - It will be also useful to review some processing techniques we've seen along the course



# JPEG: step by step


## STEP 1: Chromatic conversion

- JPEG takes in input standard RGB images
- RGB representation is redundant (there is a significant correlation among color components)
- Since we would like to remove redundancy as much as possible, JPEG converts the image in YCbCr format
- In YCbCr (similar to YUV used in old TV systems), Y is the luminance component, while Cb and Cr are chromatic components and contain less information



# JPEG: step by step

- Chromatic components downsampling
  - Our visual system is less sensitive to Cb and Cr spatial variations, then we can downsample them
  - Both images are downsampled by a factor 2 along x and y
  - For a  $N \times M$  input image, we obtain 3 images: Y ( $N \times M$ ), Cb, Cr ( $N/2, M/2$ )
- Then, the image is divided into  $8 \times 8$  blocks.
  - Given the different resolution, for every 4 Y blocks we will have 1 Cb and 1 Cr block (this represents a  $16 \times 16$  image block)
  - This configuration is called **4:1:1 scheme**



# JPEG: step by step

## ➤ DCT Transform

- Every 8x8 block (independently if Y or CbCr) is transformed by DCT
- The transform is calculated in integer form using a look-up table
- The upper left coefficient (DC) will be proportional to the mean value
- The remaining coefficients will contain increasing frequencies and will typically show sharply decreasing values

$$f(j,k) = \begin{bmatrix} 139 & 144 & 149 & 153 & 155 & 155 & 155 & 155 \\ 144 & 151 & 153 & 156 & 159 & 156 & 156 & 156 \\ 150 & 155 & 160 & 163 & 158 & 156 & 156 & 156 \\ 159 & 161 & 162 & 160 & 160 & 159 & 159 & 159 \\ 159 & 160 & 161 & 161 & 162 & 155 & 155 & 155 \\ 161 & 161 & 161 & 161 & 160 & 157 & 157 & 157 \\ 162 & 162 & 161 & 163 & 162 & 157 & 157 & 157 \\ 162 & 162 & 161 & 161 & 163 & 158 & 158 & 158 \end{bmatrix}$$

pixel-domain values



$$F(u,v) = \begin{bmatrix} 1260 & -1 & -12 & -5 & 2 & -2 & -3 & 1 \\ -23 & -17 & -6 & -3 & -3 & 0 & 0 & -1 \\ -11 & -9 & -2 & 2 & 0 & -1 & -1 & 0 \\ -7 & -2 & 0 & 1 & 1 & 0 & 0 & 0 \\ -1 & -1 & 1 & 2 & 0 & -1 & 1 & 1 \\ 2 & 0 & 2 & 0 & -1 & 1 & 1 & -1 \\ -1 & 0 & 0 & -1 & 0 & 2 & 1 & -1 \\ -3 & 2 & -4 & -2 & 2 & 1 & -1 & 0 \end{bmatrix}$$

DCT-domain values

# JPEG: step by step

- Quantization of coefficients
  - To exploit the energy compaction properties of the DCT, we need to select significant coefficients
  - JPEG does not use either zonal or threshold coding, but rather a combination of the two
- Coefficients are integer-divided by a pre-defined matrix  $Q$ , whose values are arranged according to the coefficient position
  - Before dividing, matrix  $Q$  is multiplied by a factor  $q$ , which determines the impact of the operation and then the compression
  - Most coefficients, especially in the high-frequency area, are cut. The remaining ones are strongly quantized

$$F_Q(u, v) = \left\lfloor \frac{F(u, v)}{q \cdot Q(u, v)} \right\rfloor$$


# JPEG: step by step

- The standard Q matrix is the following:  $Q(u,v)=$

$$\begin{bmatrix} 16 & 11 & 10 & 16 & 24 & 40 & 51 & 61 \\ 12 & 12 & 14 & 19 & 26 & 58 & 60 & 55 \\ 14 & 13 & 16 & 24 & 40 & 57 & 69 & 56 \\ 14 & 17 & 22 & 29 & 51 & 87 & 80 & 62 \\ 18 & 22 & 37 & 56 & 68 & 109 & 103 & 77 \\ 24 & 35 & 55 & 64 & 81 & 104 & 113 & 92 \\ 49 & 64 & 78 & 87 & 103 & 121 & 120 & 101 \\ 72 & 92 & 95 & 98 & 112 & 100 & 103 & 99 \end{bmatrix}$$

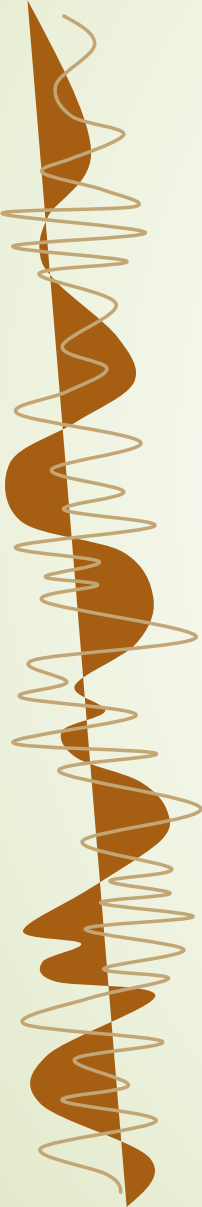
- It has been defined as to produce effective compression, while at the same time preserving the perceptual quality of the image
  - Coefficients have been refined using panels of users
  - It may be customized by the specific implementation

$$F(u,v)=\begin{bmatrix} 1260 & -1 & -12 & -5 & 2 & -2 & -3 & 1 \\ -23 & -17 & -6 & -3 & -3 & 0 & 0 & -1 \\ -11 & -9 & -2 & 2 & 0 & -1 & -1 & 0 \\ -7 & -2 & 0 & 1 & 1 & 0 & 0 & 0 \\ -1 & -1 & 1 & 2 & 0 & -1 & 1 & 1 \\ 2 & 0 & 2 & 0 & -1 & 1 & 1 & -1 \\ -1 & 0 & 0 & -1 & 0 & 2 & 1 & -1 \\ -3 & 2 & -4 & -2 & 2 & 1 & -1 & 0 \end{bmatrix}$$



$$F_Q(u,v)=\begin{bmatrix} 79 & 0 & -1 & 0 & 0 & 0 & 0 & 0 \\ -2 & -1 & 0 & 0 & 0 & 0 & 0 & 0 \\ -1 & -1 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \end{bmatrix}$$

# JPEG: step by step

- 
- After quantization:
    - The DC coefficient has usually larger values, and it is the only one to maintain a correlation across blocks
    - The AC coefficients are very sparse and usually located near to the upper left corner
  - The coding block treats differently DC and AC coefficients:
    - The DC coefficient is coded in DPCM with the previous one to exploit residual correlation
    - AC coefficients are zig-zag scanned and coded with special keywords using RLE and coefficient categories (we'll see in detail)

*NB. As you can see, JPEG is indeed a combination of different methods: it includes transform coding, RLE, DPCM and more*



# JPEG: step by step

- Zig-zag coding is used to order the coefficients in such a way that non-null values have higher probability to be scanned first, thus creating long sequences of zeros

$$F_Q(u, v) = \begin{bmatrix} 79 & 0 & -1 & 0 & 0 & 0 & 0 & 0 \\ -2 & -1 & 0 & 0 & 0 & 0 & 0 & 0 \\ -1 & -1 & 0 & \dots & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \end{bmatrix}$$

Scanned sequence: 0, -2, -1, -1, -1, 0, 0, -1, 0...

# JPEG: step by step

- Each coefficient is associated to a 8 bit codeword 'NNNNSSSS'
  - the 4 bits **NNNN** encode the previous run of zeros in RLE (this allows representing runs from 0 to 15 in length)
  - the 4 bits **SSSS** encode the coefficient range (see table below): only 10 categories are needed as the maximum value is +/-1024

Category	Range of AC value
1	-1, +1
2	-3,-2,+2,+3
3	-7,...,-4,+4,...,+7
4	-15,...,-8,+8,...,+15
...	...
10	-1023,...,-512,+512,...,+1023

*NB. The specific value within the category is sent separately. The  $i$ -th category requires exactly  $i$  bits for indexing (**precision bits**)*

# JPEG: step by step

- The number of codewords used is  $16 \times 10 = 160$
- Part of the remaining codewords are used for special cases:
  - **Continuation code:** if a run is longer than 15, a run continuation code is sent with the word 11110000 (category 0 does not exist!)
  - **End of block:** after the last non-null coefficient, and EOB code is sent to indicate that nothing will follow, with codeword 00000000

In our example, the scanned sequence was: 0, -2, -1, -1, -1, 0, 0, -1, 0s

Then we will have:

run 1, category 2: 00010010

run 0, category 1: 00000001

run 0, category 1: 00000001

run 0, category 1: 00000001

run 2, category 1: 00100001

EOB: 00000000

+ precision bits: 01, 0, 0, 0, 0

# JPEG Step by step

- Finally, some redundancy still remains in the codewords
  - Some of them are much more frequent than other
  - For instance, EOB is always present, lower categories and shorter runs are more probable, etc.
- To exploit this redundancy we can use an entropy coder
  - the standard uses an Huffman LUT
  - alternatives are also possible

In our example we will have:

11100101, 000, 000, 000, 110110, 1010

+ precision bits 01, 0, 0, 0, 0

Zero	Cat.	Code	Codeword
0	1	2	00
0	2	2	01
0	3	3	100
0	4	4	1011
0	5	5	11010
0	6	6	111000
0	7	7	1111000
1	1	4	1100
1	2	6	111001
1	3	7	1111001
1	4	9	111110110
2	1	5	11011
2	2	8	11111000
3	1	6	111010
3	2	9	111110111
4	1	6	111011
5	2	7	1111010
6	1	7	1111011
7	1	8	11111001
8	1	8	11111010
9	1	9	111111000
10	1	9	111111001
11	1	9	111111010
EOB		4	1010

# JPEG: Decoding

- Decoding will follow the same processing chain in inverse order (the technique is symmetric)
  - Apply Huffmans decoding to reconstruct NNNNSSSS codewords
  - Reconstruct the sequence and values of coefficients from NNNNSSSS codewords + precision bits, and place it into the block following the zig-zag order
  - De-quantize the block multiplying by the  $q \cdot Q$  matrix
  - Apply the inverse DCT transform to reconstruct pixel values
  - Upsample Cb and Cr components
  - Apply the color transform  $YCbCr \rightarrow RGB$
- The decoded image will be different from the original
  - The two lossy (irreversible) operations are CbCr downsampling and especially, coefficients quantization

# JPEG: Performance

- To compression factor is highly variable and depends indirectly on the  $q$  parameter
  - In JPEG we cannot set the compression, we set the desired quality. The resulting compression depends on the image content
- The code contains an header (negligible), the DPCM bits, the Huffman code bits, and the precision bits
  - In our example we have (for 1 image block):
    - DPCM: depends on correlation with previous block, let's use an average of 4 bits
    - Huffman: 111001010000000001101101010, total 27 bits
    - Precision: 010000: 6 bits
  - Total: 37 bits, instead of  $8 \times 8 \times 8 = 512$  bits
    - Compression factor  $C_f = 512/37 = 13.8$
    - Bitrate  $r_b = 8/13.8 = 0.58$  bpp



# JPEG Performance

- The result of decoding in our sample block will be:

$$\hat{f}(j,k) = \begin{bmatrix} 144 & 146 & 149 & 152 & 154 & 156 & 156 & 156 \\ 148 & 150 & 152 & 154 & 156 & 156 & 156 & 156 \\ 155 & 156 & 157 & 158 & 158 & 157 & 156 & 155 \\ 160 & 161 & 161 & 162 & 161 & 159 & 157 & 155 \\ 163 & 163 & 164 & 163 & 162 & 160 & 158 & 156 \\ 163 & 163 & 164 & 164 & 162 & 160 & 158 & 157 \\ 160 & 161 & 162 & 162 & 162 & 161 & 159 & 158 \\ 158 & 159 & 161 & 161 & 162 & 161 & 159 & 158 \end{bmatrix}$$

- Subtracting it from the original we obtain:

$$e(j,k) = \begin{bmatrix} -5 & -2 & 0 & 1 & 1 & -1 & -1 & -1 \\ -4 & 1 & 1 & 2 & 3 & 0 & 0 & 0 \\ -5 & -1 & 3 & 5 & 0 & -1 & 0 & 1 \\ -1 & 0 & 1 & -2 & -1 & 0 & 2 & 4 \\ -4 & -3 & -3 & -1 & 0 & -5 & -3 & -1 \\ -2 & -2 & -3 & -3 & -2 & -3 & -1 & 0 \\ 2 & 1 & -1 & 1 & 0 & -4 & -2 & -1 \\ 4 & 3 & 0 & 0 & 1 & -3 & -1 & 0 \end{bmatrix}$$



MSE = 5.1  
PSNR = 41 dB

# JPEG: Examples

- JPEG coding of Lena at maximum quality (100)



original



Cf = 3, PSNR=37.50 dB

# JPEG: Examples

- JPEG coding of Lena at medium and low quality (50, 10)



$C_f = 22$ , PSNR=29.81 dB




$C_f = 55$ , PSNR=25.28 dB

NB. **tiling** and **ringing artifacts** appear, due to block transform coding!

# JPEG2000

- JPEG2K is an evolution of JPEG standardized almost 10 year later
  - At that time, the main technological breakthrough in this field was the Wavelet Transform (DWT)
  - DWT is pretty different from DFT and DCT, in that it produces a transformed signal which is in a mixed **space-frequency domain**
  - Besides this, the wavelet domain is said to be better adapted to the characteristics of the human visual system
- This led to the development of a rather different approach of transform coding
  - Not only a different transform is used, with better performances, but also a number of **additional functionalities** are introduced (e.g., scalability, resilience, progressive reconstruction, regions of interest, ...)

# JPEG2000 and Wavelets

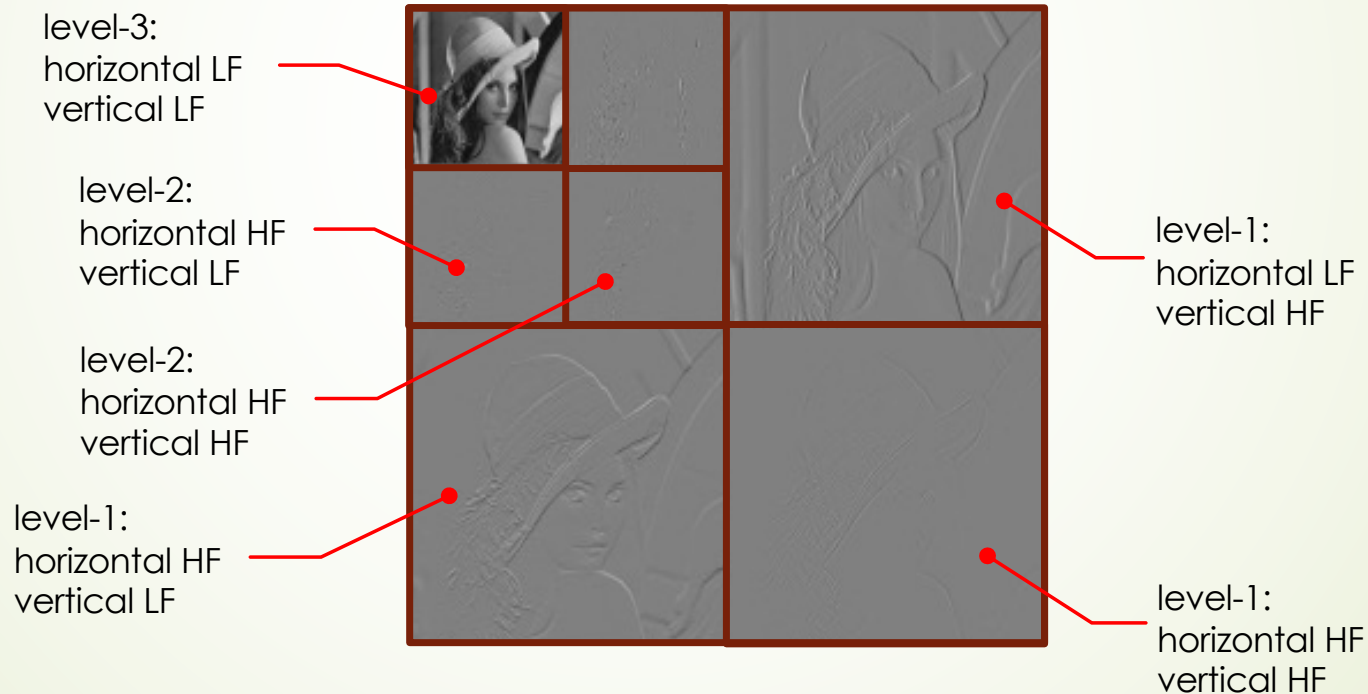


- A complete study of wavelets is out of the scope of this course
- A very simple introduction of the concept is the following:
  - We want to detect the presence of some frequency content in a given area of the image
  - We can do that by analyzing the image with appropriate functions (wavelets) containing localized frequencies at various intensities
  - Wavelets can be generated by scaling and translating a unique function, called **wavelet mother function**



# JPEG2000 and Wavelets


- In practical terms, what we do is to use special filters and to analyze the image at different scales factors
  - What we get is a kind of **pyramidal representation**






# JPEG2000: Scheme

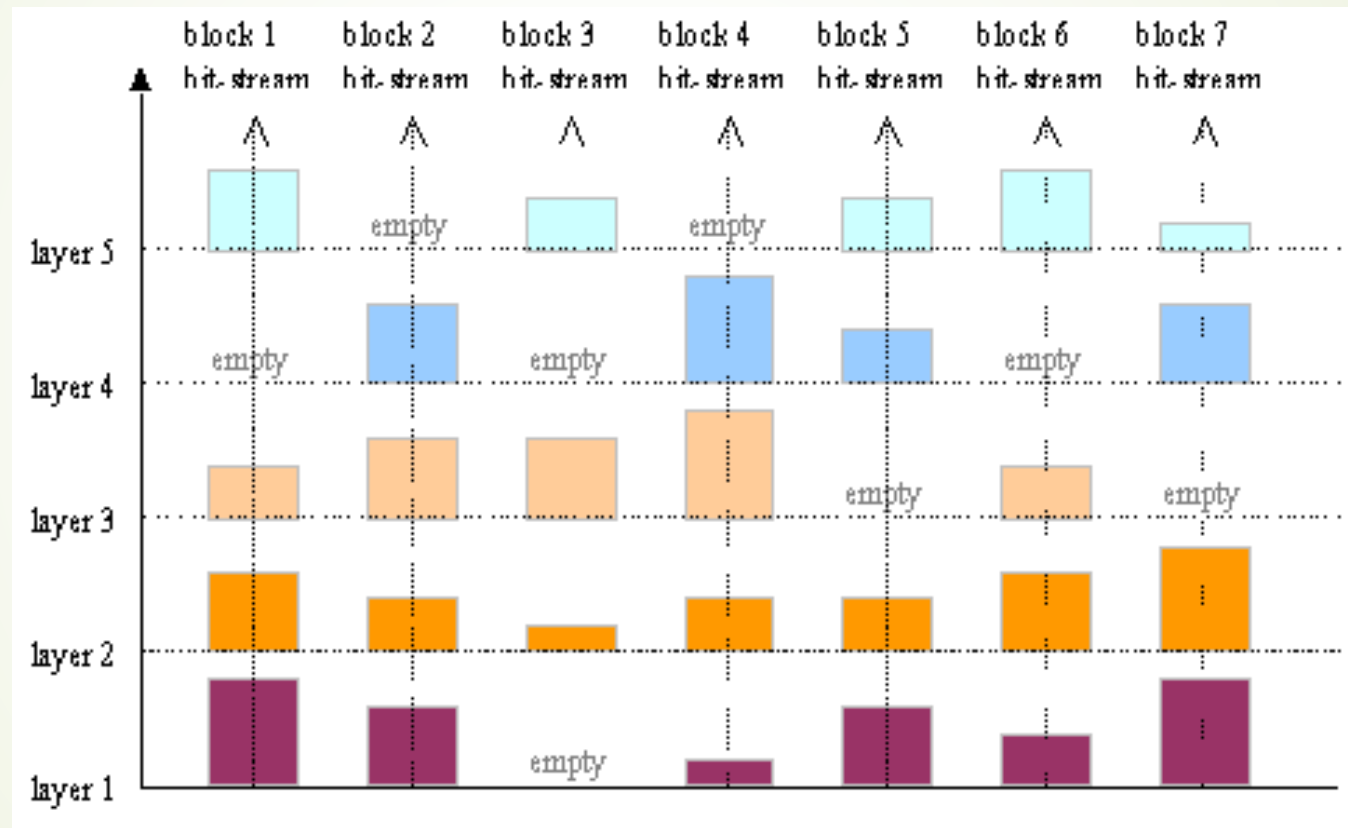
- The encoder is apparently quite similar to JPEG
  - The biggest difference is the encoder (EBCOT: embedded block coding with optimized truncation)
  - The transform is applied to blocks within each sub-band of the pyramidal structure
  - The bits associated to the various blocks have different importance according to the local content: an optimized allocation of bits allows maximizing the rate-distortion function for any given stream size



# JPEG2000: Block coding



# JPEG2000: Bit allocation



# JPEG2000: Scalability and Rols

- Thanks to the code structure we can transmit an image at different resolution or quality in a progressive way
- We can also transmit an image with differentiated quality (better quality in the area of interest)



Original with Rol



Reconstructed

# JPEG2000: Performance

- JPEG2000 provides better visual and SNR quality wrt JPEG and allows to achieve much higher compression factors



JPEG at  $C_f$  55



JPEG2000 at  $C_f$  55



## What we've learned in this section

- Images are huge data objects, they need compression
- Images can be compressed because they are redundant
- Coding may introduce degradation or not. If it does, we have to measure the image quality
- Quality is related to our (human) perception
- Coding without losing information has strong limitations
- Lossy coding allows larger compressions at the price of some image degradation (it's a trade-off)
- Current standards are based on sophisticated implementations of block transform coders
- Even if it has almost 30 years and there are better performing techniques, JPEG is still the most diffused image coding standard