




FUNDAMENTALS OF IMAGE AND VIDEO PROCESSING

Part 5: Middle-level image processing

Middle-level image processing

- In this section we introduce intermediate-level operations, whose purpose is to **convert the image into more meaningful** (semantically richer) **representations**
 - For instance, starting from a pixel-level representation we want to describe the image in terms of contained “objects”
 - Each object could be described by its shape, color, texture, etc.
 - At this level, we are still not able to define what an object is, but just its **appearance**.
- Middle-level processing produces descriptions that are easier to use for high-level systems
 - Based on descriptors, high-level could go further in semantic interpretation, trying to make sense of the image content

Middle-level processing: example



- Starting from the image on the left, we want to describe the objects contained in the scene (right)
 - We could, e.g., extract areas that are uniform according to some properties, and describe them in terms of shape and texture
 - For instance, the blue 'blob' on the left could be described as compact shape with sharp angles, mostly convex with a major concavity on the bottom, with a dominant dark-red color.
 - We are still unable to associate such blob to the concept of 'car'

Image source: A. Kirillov et al., CVPR 2019

From middle- to high-level

- ▶ Later on, the blobs could be associated to visual concepts
 - ▶ This requires sophisticated machine learning techniques




Image source: towardsdatascience.com

Image descriptors


- ▶ We'll introduce four main types of descriptors:
 - ▶ **Contours:** boundaries of objects
 - ▶ **Regions:** shapes associated to objects
 - ▶ **Textures:** arrangement of colors on the object surface
 - ▶ **Structures:** groups of objects with spatial relationships
- ▶ The four descriptors are somewhat complementary
 - ▶ Contours enclose regions
 - ▶ Regions contain textures
 - ▶ Groups of regions/contours create structures
- ▶ The techniques used to extract such descriptors from images, however, are quite different

The process of extracting image descriptors

- We distinguish 3 main phases:
 - **Detection:** it is the action of revealing the desired descriptor from the image (e.g., detecting edge points to produce an edge-map)
 - **Representation:** it is the action of associating an appropriate description to each detected primitive (e.g., representing a chain of edge point as a 2D curve, or contour)
 - **Feature extraction:** it is the action of associating a set of quantitative parameters to each detected primitive (e.g., representing a contour in terms of length, shape, closeness, curvature, etc.)
- Each phase can be implemented in different ways (algorithms), with relevant pros and cons

Contours

- The first descriptors we introduce are image contours
- The **expressivity of contours** in representing the nature of objects is rather evident



Contour extraction process



- ▶ As we've seen, we distinguish 3 phases:
 - ▶ **Edge detection:** we convert an image (grayscale or color) into a binary map of edges
 - ▶ **Contour extraction:** we scan the edge points that have been detected at the previous step to create connected chains of contours (curves)
 - ▶ **Feature extraction:** we associate a set of parameters to each contour chain

Contour extraction: example





original image

detection



edge map

extraction



contour

features




Contour 1

- length: 1252 pels
- Closure: closed
- Shape: polygonal
- Corners: 5
- ...
- ...

description

Edge detection

- Edges are characterized by steep luminance/color variations that are present in an image in the presence of object borders
 - Since edges are associated to steep variations of the image function, they can be detected by analyzing spatial high-frequencies
 - Typical approaches are therefore based on the use of gradient filters



Sobel edge detector: example

x derivative



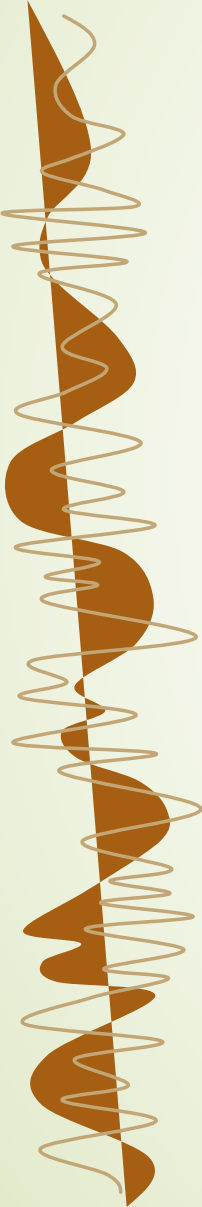
y derivative



modulus



Sobel pseudo-code



```
int img[N][M], edges[N][M]           // input, output images
int grad_x[N][M], grad_y[N][M]       // directional gradient images
int kernel_x = [[1,2,1],[0,0,0],[-1,-2,-1]] // Sobel FIR x direction
int kernel_y = [[1,0,-1],[2,0,-2],[1,0,-1]] // Sobel FIR y direction

load (img)
grad_x = convolve(img, kernel_x)      // FIR filtering with kernel x
grad_y = convolve(img, kernel_y)      // FIR filtering with kernel y
for n in 0...N-1 {                    // raster scan image
    for m in 0...M-1 {                // " "
        tmp = (grad_x[n][m]^2+grad_y[n][m]^2) // square mod
        if (tmp > THR) edges[n][m] = 1      // thresholding
        else edges[n][m] = 0                // " "
    }
}
```

Gradient operators: pros and cons

■ PROS

- Easy to implement
- Limited complexity (two 3x3 FIR filters)
- Relatively accurate if image is not too noisy

■ CONS

- Hardly detect weak contours
- Disconnected and thick contours (depends on thresholding)
- Imprecise localization

Canny edge detector

- John Canny proposed to address the problem of optimal edge detection in a formal way
- An “optimal” detector should provide:
 - low misdetection (catch as many real edges as possible)
 - good localization (identify the center of the edge)
 - low rate of false edges (edges should be marked only once and noise should not create false edges)
- To this end Canny proposed a method based on the optimization of a functional, defined as the sum of 4 exponential terms
 - In actual implementations, the above optimization is approximated by a more traditional sequence of filters

Canny: typical implementation

- The detector follows 5 steps:
 - Apply a **Gaussian filter** to remove noise (typically, a 5x5 FIR kernel)
 - Calculate the **gradients along x and y** (similar to Sobel), and associate to each point an intensity and an angle (quantized to 4 directions)
 - Perform **non-maximum suppression**, a kind of thinning where the edge strength is compared to the neighbors along the edge direction and only the point with larger gradient is selected
 - Apply **double thresholding**: pixels are marked as strong edges, weak edges or suppressed according to a lower and an upper threshold
 - Edge **tracking by hysteresis**: weak edges are preserved only if they are 8-connected to at least a strong one

Canny edge detection: example



Canny detector: pros and cons


➤ PROS

- Very good localization
- Thin contours
- High-continuity of lines, also for weak contours
- Highly resistant to noise

➤ CONS

- Several parameters/thresholds
- Relatively complex

Contour representation & description



- Edge detectors produce in output binary images
- This is not yet a true contour representation
 - We still need to process the edge image and extract the contour chains, in the form of 2D curves.
- To this purpose, we need appropriate algorithms to follow sequences of connected contour points
- The most known algorithm is the so-called **Freeman chain code**

Textures

- The last descriptor we introduce is the texture
- Textures complement contours and shapes in revealing (or hiding) the nature of an object



Textures

- A texture is the visual appearance of a surface. It can be thought as a spatial arrangement of colors (pattern) showing some kind of regularity
 - In the simplest case, it can be a uniform color
- It is a complementary feature that, associated to a shape (contour or region), completes the description of an object



3D shape



shape + texture

Examples of textures




Source: Brodatz texture database

Texture analysis


- We typically don't "detect" textures, but we rather analyze them to:
 - define parameters that can identify a given texture (e.g., to evaluate texture homogeneity in segmentation)
 - detect texture irregularities (e.g., finding anomalies in textured surfaces for visual inspection purposes)
 - classify textures (e.g., in object detection, when the texture is a characteristic feature of a given object)
 - define models to synthesize similar textures (e.g., in computer graphics and virtual reality)

Texture analysis: examples

- Using texture homogeneity feature in segmentation




Artificial test image




Segmentation based
on texture similarity

Texture analysis: examples

► Texture fault detection



A sample of tissue
with a defect




Defect detection basd
on texture anomaly

Texture analysis: example

► Texture synthesis




Texture sample



Synthesized pattern with
texture sample properties

Transform-based methods

- Transforms (e.g., DFT but not only) provide information about the energy distribution of an image in the frequency domain
 - Looking at transform coefficients I can perceive if an image contains higher or lower frequencies, as well as their dominant directions
 - This is very much related to the textures present in the image, in fact, textures produce frequency peaks in specific zones of the transform



texture sample

transformed sample

Source: G. Dougherty, M.A. Haidekker, "Medical Image Processing: Techniques and Applications", Springer, 2011